Lemma — Pre-release Brief

Hey everyone 👋

Before the hackathon opens fully, here's a little context on what Lemma is, why we're building it, and how to start shaping your ideas.

What is Lemma?

Lemma is an open-source workspace where humans and AI agents work together.

Not just chatting with an AI. Not just asking an agent to generate something once.

Lemma gives agents a place to actually *do work* — with tasks, tables, files, approvals, permissions, and apps that humans use too.

Think of it as a shared workspace where:

- humans assign, review, and approve work
- AI agents own tasks and update state
- everything lives in structured tables and files
- work doesn't disappear inside a chat thread

Why does this exist?

Most AI demos today end in a chat log. That's fine for simple questions — but real work is different.

Real work has owners, status, deadlines, approvals, shared context, history, and rules about who can do what. A chat thread isn't built for any of that.

If an AI agent qualifies a lead, reviews a PR, drafts an email, or triages bugs, the output shouldn't be another paragraph in chat. It should land somewhere useful:

- a row in a table
- a task in a queue
- a draft waiting for approval
- a workflow that picks up again tomorrow

> The best AI products already work this way. **Cursor** lands its work as diffs in your editor. **Gamma** turns a prompt into an editable deck. **Granola** turns a meeting into

> structured notes. Same shape every time — an agent works in the background, and structured output lands in a real UI. **Lemma is that shape for your work.**

The simple idea

> **Chat is where work starts. Lemma is where work lives.**

You can still use the chat tools you already use — Slack, WhatsApp, Telegram, Gmail, Teams. But underneath, Lemma holds the real state of the work.

What is a pod?

A **pod** is a self-contained workspace for one process or team. It's just a directory of plain files, so you can export it, edit it, and share it.

Inside a pod	What it gives you
Tables	Structured data — leads, bugs, tasks, invoices, applications
Files	Markdown knowledge — playbooks, notes, rules, preferences
Agents	AI workers with a role and *scoped* permissions
Workflows	Steps combining agents, logic, waits, and approvals
Approvals	Points where a human must review before something happens
Apps (desks)	Interfaces humans use to operate the system
Surfaces	Chat/email channels wired to the same workspace

Instead of a random AI chatbot, you build a working *system*.

A few things you'll like

- **It's yours.** Open-source, runs on your own machine. Your data stays where you put it.
- **It uses the subscription you already pay for.** Agents can run on your Claude or ChatGPT login — no separate API bill — or any key/model you want.

- **Your coding agent can build the pod for you.** Because a pod is just files, you can describe the system to Claude Code / Cursor / Codex, let it write the pod, and import it. The same agent can then test it.

Why this matters for the hackathon

We don't want you to build "just another wrapper around ChatGPT." Think bigger: build something where an AI agent actually *participates* in a real workflow.

A few directions:

- an AI bug-triage desk for a dev team
- a PR review + release-readiness assistant
- a job-application command centre
- a customer support desk
- a meeting-to-execution operator
- a personal life command centre
- a founder follow-up CRM
- a learning companion that remembers your progress

The best projects won't just generate text. They'll show how work moves: **input** → **state** → **action** → **approval** → **outcome.**

Start thinking about your workflow

Pick a real one, and ask:

1. What information needs to be tracked?
2. Who owns each step?
3. Where should an agent help?
4. Where must a human approve?
5. What should be visible in a table or app?
6. What should happen automatically?
7. What should *never* happen without permission?

That's the Lemma way of thinking.

The bet underneath it all

Coding agents can now build almost any one-off app — which means the world is about to drown in unmaintainable, agent-generated software. Lemma gives agents a shared framework to build with, and a place to run what they build, with humans and approvals around them.

The unlock isn't more agent autonomy. It's *approvals* — agents doing real work, with a human in the loop where it counts.

Access

Lemma is still pre-release, so full access isn't open yet. For now, use this brief to understand the direction and start shaping ideas. When access opens, you'll build on Lemma as an open-source workspace for humans and AI agents working together.

One line to remember:

> **Don't build a chatbot. Build a workspace where AI can do real work.**